

Terraform Task

Name: Shaik Khaja Basha

Batch : Batch 11

Date : 24.07.2025

Task : ignore changes & replace triggered

1. Explore the life cycle of ignore_changes and replace_triggered_by ?

Ans: In Terraform, ignore_changes and replace_triggered_by are lifecycle meta-arguments that control how resources are managed. ignore_changes prevents Terraform from attempting to update specific attributes that are being modified outside of Terraform, while replace_triggered_by forces a resource to be replaced when specific referenced resources or variables change

ignore_changes:

purpose:

- To ignore changes to specific attributes of a resource that are being managed outside of Terraform.
- Even if instance_type changes in code or tfvars, Terraform will **not trigger a replacement** or update the resource.

Use Cases:

- When an external process (like an Azure Policy) modifies resource attributes, and Terraform should not attempt to revert those changes.
- When dealing with dynamic data that may change after resource creation, but should not trigger a resource replacement.

Example:

Ignoring changes to tags on an AWS instance that are being managed by an external tagging process.

working:

Terraform will still manage the resource's creation and destruction, but it will not attempt to reconcile the specified attributes with the Terraform configuration.

Syntax:

```
resource "aws_instance" "example" {
    ami      = "ami-123456"
    instance_type = var.instance_type

    lifecycle {
        ignore_changes = [instance_type]
    }
}
```

Scenario:

You manage an Azure App Service. Operations team might change settings manually, and you don't want Terraform to revert their changes every time.

Example:

```
lifecycle {  
  ignore_changes = [ app_settings ]  
}
```

Terraform will not re-apply changes if only app_settings are different in the state vs config.

Effect:

- Terraform won't recreate the resource if only the timestamp changes
- Useful when you want to keep the resource stable despite volatile attributes

replace_triggered_by:**Purpose:**

To force a resource to be recreated when specific referenced resources or variables change

Use Cases:

- When a resource depends on another resource whose changes should trigger a full replacement of the dependent resource.
- When a resource needs to be replaced when a specific variable or attribute of another resource changes.
- If the local_file.example resource is modified or recreated, null_resource.example will also be replaced.

Example:

Replacing a Google Compute Instance when its startup script changes.

How it works:

When the referenced resource or variable changes, Terraform will plan a replacement of the resource, effectively destroying and recreating it.

Syntax:

```
resource "null_resource" "example" {  
  triggers = {  
    value = var.value  
  }  
  
  lifecycle {  
    replace_triggered_by = [  
      var.value,  
      local_file.example    # or another resource  
    ]  
  }  
}
```

Example with Random Provider:

```
resource "random_id" "seed" {  
  byte_length = 4  
}  
  
resource "null_resource" "dependent" {  
  triggers = {  
    seed = random_id.seed.hex  
  }  
  
  lifecycle {  
    replace_triggered_by = [random_id.seed] # Recreate when seed changes  
  }  
}
```

Effect:

- null_resource.dependent will be destroyed/recreated whenever random_id.seed changes
- Creates an explicit dependency link

Difference between ignore_changes and replace_triggered_by

ignore_changes	replace_triggered_by
ignore_changes prevents updates to specific attributes	replace_triggered_by forces a full resource replacement.
ignore_changes is useful for managing resources where external processes make changes	replace_triggered_by is useful for managing resources that depend on other resources that may change.
ignore_changes allows for shared management of a resource between Terraform and external processes	replace_triggered_by ensures consistency by replacing the resource when a specific dependency changes